

# Lecture 11 - Contents

An overview of the main sections in this lecture.

## Part 1

Medical Reward Modeling

## Part 2

Policy Optimization

## Part 3

RLHF in Clinical Practice

## Hands-on

RLHF Pipeline

This outline is for guidance. Navigate the slides with the left/right arrow keys.



Lecture 11:

# **RLHF in Healthcare: Aligning AI with Medical Expertise**

**Ho-min Park**

[homin.park@ghent.ac.kr](mailto:homin.park@ghent.ac.kr)

[powersimmani@gmail.com](mailto:powersimmani@gmail.com)

# RLHF Overview in Healthcare

## What is RLHF?

Reinforcement Learning from Human Feedback aligns AI models with medical expertise through iterative feedback loops, ensuring outputs match clinical standards and safety requirements.



### Clinical Alignment

Ensures AI decisions match medical standards

### Safety Enhancement

Reduces harmful or inappropriate outputs

### Continuous Improvement

Adapts to evolving medical knowledge

### Expert Integration

Incorporates physician judgment directly

**Part 1/3:**

# **Medical Reward Modeling**

- 1.** Clinical Preference Learning
- 2.** Expert Feedback Collection
- 3.** Annotation Interfaces
- 4.** Preference Dataset Creation
- 5.** Reward Model Architecture
- 6.** Bradley-Terry Model
- 7.** Uncertainty Estimation

# Clinical Preference Learning

## What is Clinical Preference Learning?

Process of capturing expert medical judgment through pairwise comparisons of model outputs, enabling the AI to learn what constitutes high-quality clinical responses.

## Collection Methods

- Pairwise Comparisons: Experts choose between two model outputs
- Ranking Tasks: Order multiple outputs by quality
- Absolute Scoring: Rate individual outputs on fixed scales
- Natural Language Feedback: Detailed written critiques

## Key Considerations

- Inter-annotator Agreement: Ensure consistency across experts
- Sample Diversity: Cover wide range of clinical scenarios
- Quality Control: Regular calibration sessions
- Expert Qualifications: Board-certified specialists in relevant domains

# Expert Feedback Collection

## Annotation Protocol

Structured guidelines ensure consistent, high-quality feedback from medical experts across all evaluation tasks.

## Collection Pipeline



5

### Quality Monitoring

Track agreement metrics and outliers

## Quality Assurance

- Gold Standard Examples: Pre-annotated cases for validation
- Inter-rater Reliability: Cohen's Kappa, Fleiss' Kappa
- Regular Audits: Review annotation quality periodically
- Feedback Loops: Discuss disagreements and edge cases

# Annotation Interfaces

## Interface Design Principles

User-friendly annotation tools maximize expert efficiency and reduce cognitive load during evaluation tasks.

## Key Features

- Side-by-Side Comparison: View outputs simultaneously
- Contextual Information: Patient history, relevant guidelines
- Quick Actions: Keyboard shortcuts for common tasks
- Progress Tracking: Visual indicators of completion
- Comment System: Add detailed notes and rationale

## Workflow Optimization

- Adaptive Sampling: Focus on uncertain cases
- Break Reminders: Prevent annotation fatigue
- Session Management: Save progress automatically
- Mobile Compatibility: Annotate on various devices

# Preference Dataset Creation

## Dataset Structure

Preference datasets contain pairs of model outputs with expert rankings, forming the foundation for reward model training.

## Data Components



### 1. Input Prompt

*Clinical query or task description*



### 2. Output A

*"Patient presents with fever and cough. Recommend chest X-ray and CBC."*



### 3. Output B

*"Recommend immediate chest X-ray, CBC, and respiratory pathogen panel given symptoms."*



### 4. Preference Label

Output A



Output B ✓

*Expert prefers B: More comprehensive diagnostic approach*



### 5. Confidence Score (Optional)

Strength: High (0.9) - Clear preference for comprehensive testing



## 6. Rationale

Expert's reasoning: "Output B includes respiratory pathogen panel which is essential for differential diagnosis in current clinical context."

### Dataset Statistics

**10K-100K+**

Preference Pairs

**50/50**

Balance (A vs B)

**15+ Specialties**

Clinical Coverage

**85-95%**

Agreement Rate

# Reward Model Architecture

## Model Structure

Reward models predict scalar scores for model outputs, learned from expert preference data to guide policy optimization.

## Architecture Components

### Input Layer

Tokenized Output Pairs (Output A, Output B)



### Base Encoder (Transformer)

BERT, GPT, or Medical Domain Transformer



### Pooling Layer

Aggregate Sequence Representations (Mean/Max/CLS)



### Reward Head (Linear Layer)



## Loss Function

Bradley-Terry / Ranking Loss

### Training Process

#### 1 Input

Pairs of outputs with preference labels

#### 3 Loss Calculation

Penalize incorrect rankings

#### 2 Forward Pass

Compute reward scores for both outputs

#### 4 Optimization

Adam optimizer with LR scheduling

# Bradley-Terry Model

## Mathematical Foundation

The Bradley-Terry model converts reward scores into probabilities for pairwise comparisons, providing a principled approach to preference learning.

## Model Formula & Visualization

$$P(A > B) = \sigma(r(A) - r(B)) = 1 / (1 + \exp(-(r(A) - r(B))))$$

Output A

$$r(A) = 0.8$$

Output B

$$r(B) = 0.3$$

VS



## Higher Reward Difference → Stronger Preference Probability

💡  $\sigma$  (sigmoid function) ensures probability output between 0 and 1

### Training Objective

- Maximize log-likelihood of observed preferences
- Loss =  $-\log(P(\text{preferred} > \text{not\_preferred}))$
- Gradient descent updates reward model parameters
- Converges to scores matching expert preferences

# Uncertainty Estimation in Reward Models

## Why Uncertainty Matters

In medical AI, knowing when the reward model is uncertain helps identify cases requiring additional expert review or model improvement.

## Uncertainty Spectrum Visualization



High Confidence

✓ Deploy

Model is certain, safe to use



Medium Confidence

⚠ Review

Uncertain, flag for expert



Low Confidence

🚫 Block

High uncertainty, do not use

## Estimation Methods



### Ensemble Methods

Train multiple reward models with different initializations, measure prediction variance



### Bayesian Approaches

Model weight uncertainty with probability distributions (e.g., Bayesian Neural Networks)



### Monte Carlo Dropout

Apply dropout during inference for variance estimation across multiple forward passes



### Calibration

Ensure predicted confidence matches empirical accuracy using calibration techniques

## Applications

- Active Learning: Query experts on high-uncertainty cases
- Safe Deployment: Flag uncertain predictions for human review
- Model Improvement: Identify areas needing more training data
- Confidence Intervals: Provide uncertainty bounds with predictions

## **Part 2/3:**

# **Policy Optimization in Medical AI**

- 1.** PPO in Medical Applications
- 2.** Direct Preference Optimization (DPO)
- 3.** Safety Constraints
- 4.** KL Divergence Control
- 5.** Exploration vs Exploitation
- 6.** Online vs Offline RLHF

# PPO in Medical Applications

## Proximal Policy Optimization (PPO)

PPO is the most widely used algorithm for RLHF, providing stable updates that prevent catastrophic policy changes in medical AI systems.

### Key Features



#### Clipped Objective

Limits policy updates to prevent drastic changes



#### Trust Region

Maintains proximity to previous policy



#### Stability

Reduces training instability vs vanilla PG



#### Sample Efficiency

Reuses data multiple times per iteration

### Training Process (Iterative Loop)

1

2

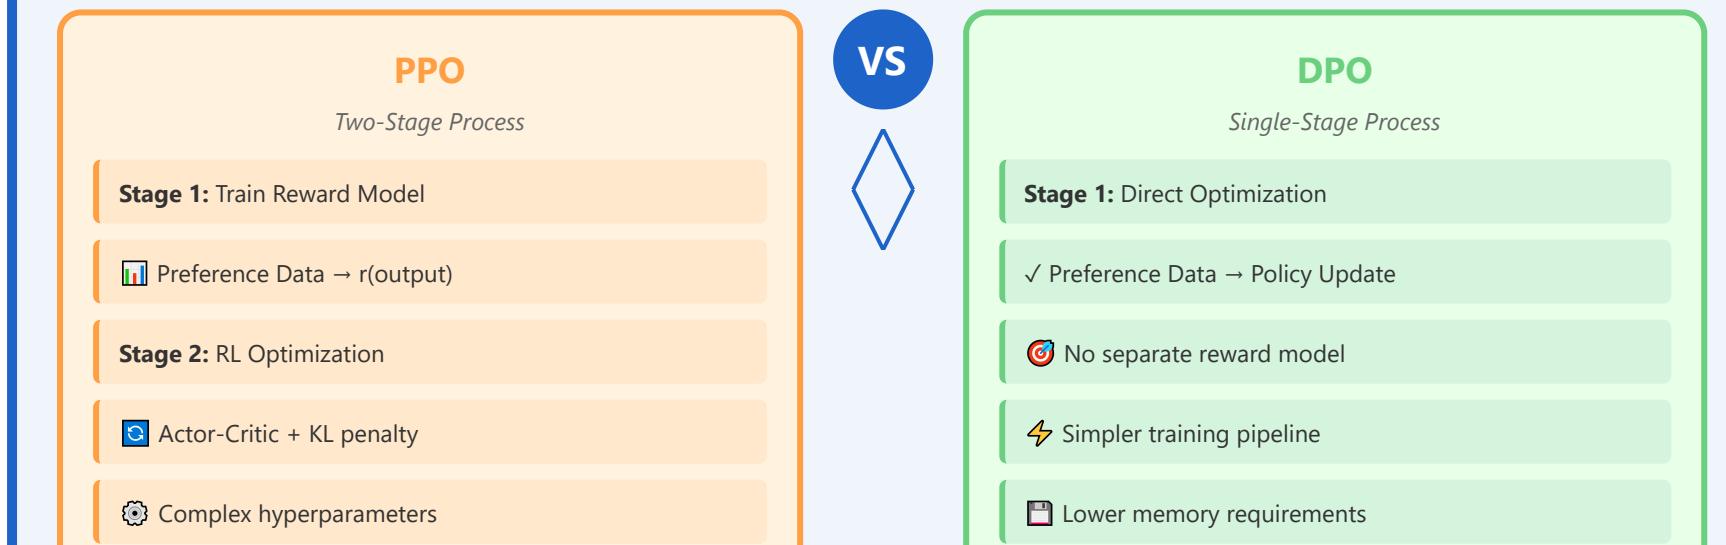
3

# Direct Preference Optimization (DPO)

## What is DPO?

DPO directly optimizes the language model using preference data, eliminating the need for a separate reward model and RL training loop.

## DPO vs PPO: Architecture Comparison



# Safety Constraints in Medical RLHF

## Why Safety Constraints?

Medical AI requires hard constraints to prevent harmful outputs that could endanger patients, regardless of reward optimization.

## Multi-Layer Safety Architecture



### Layer 1: Hard Medical Rules

Never violate established clinical guidelines (e.g., contraindications, age restrictions)



### Layer 2: Dosage Limits

Enforce safe medication dosing ranges based on patient factors (weight, age, renal function)



### Layer 3: Contraindication Checking

Prevent dangerous drug interactions and allergic reactions



### Layer 4: Appropriate Scope

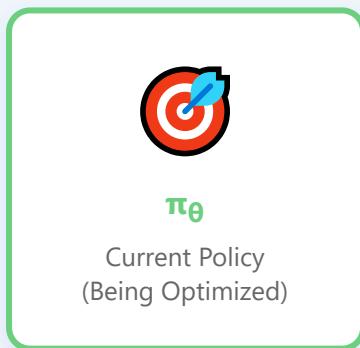
Stay within system's trained competency domain

# KL Divergence Control

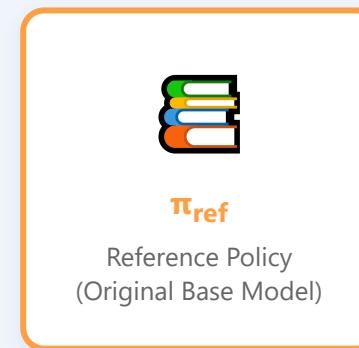
## Purpose of KL Divergence

KL divergence measures how much the optimized policy deviates from the original base model, preventing harmful divergence from foundational knowledge.

## Visual Representation



$$\text{KL}(\pi_\theta \parallel \pi_{\text{ref}})$$



‡ KL measures distribution difference



# Exploration vs Exploitation in Medical AI

## The Dilemma

Balancing exploration (trying new approaches) with exploitation (using known-good approaches) is critical in medical AI where safety is paramount.

## The Balance Concept



## Exploration Strategies



Entropy Bonus



Temperature Sampling

# Online vs Offline RLHF

## Offline RLHF

Train on fixed datasets of preferences collected in advance. More common in medical AI due to safety requirements.

## Offline Advantages

- Safety: All data reviewed before training
- Reproducibility: Fixed dataset enables consistent experiments
- Expert Efficiency: Collect preferences in batches
- Lower Risk: No live patient interaction during training

## Online RLHF

Continuously collect new preferences during deployment and retrain. Riskier but more adaptive to evolving needs.

## Online Advantages

- Adaptation: Quickly respond to new medical knowledge

## **Part 3/3:**

# **RLHF in Clinical Practice**

- 1.** Clinical Outcome Rewards
- 2.** Patient Satisfaction Metrics
- 3.** Safety-Critical RLHF
- 4.** Continuous Learning Systems
- 5.** Feedback Loop Design
- 6.** A/B Testing Framework
- 7.** Performance Monitoring
- 8.** Failure Mode Analysis
- 9.** Case Study & Hands-On

# Clinical Outcome Rewards

## Outcome-Based Reward Design

Reward models can be trained not just on preferences, but on actual clinical outcomes to align AI with real-world treatment success.

## Outcome Metrics

- Survival Rates: Long-term patient outcomes
- Readmission Rates: 30-day hospital readmissions
- Symptom Improvement: Patient-reported outcomes
- Complication Rates: Post-treatment adverse events
- Quality of Life: QALY, functional status measures

## Implementation Challenges

- Long Feedback Loops: Outcomes take time to materialize
- Confounding Factors: Many variables affect outcomes
- Sample Size: Need large datasets for statistical power
- Ethical Considerations: Can't experiment with patient health

# Patient Satisfaction Metrics

## Why Patient Satisfaction Matters

Beyond clinical accuracy, AI systems should provide positive patient experiences, which correlate with treatment adherence and outcomes.

## Satisfaction Dimensions

- Clarity: Was the explanation understandable?
- Empathy: Did the AI show appropriate compassion?
- Completeness: Were all questions answered?
- Trust: Did the patient feel confident in the advice?
- Efficiency: Was the interaction time-appropriate?

## Collection Methods

- Post-Interaction Surveys: Rating scales, open feedback
- Implicit Signals: Time spent reading, follow-up questions
- Behavioral Data: Appointment adherence, prescription fills
- Sentiment Analysis: Analyze patient language patterns

# Safety-Critical RLHF

## Medical Error Prevention

RLHF systems in healthcare must have multiple safety layers to prevent errors that could harm patients.

## Safety Mechanisms

- Pre-Deployment Testing: Extensive validation before clinical use
- Conservative Defaults: Err on side of caution when uncertain
- Human-in-the-Loop: Require expert confirmation for critical decisions
- Redundant Checks: Multiple independent safety verifications
- Fail-Safe Modes: Graceful degradation when errors detected

## Critical Failure Modes

- Medication Errors: Wrong drug, dose, or timing
- Diagnostic Misses: Failing to identify serious conditions
- Inappropriate Reassurance: Downplaying concerning symptoms
- Scope Violations: Advising beyond system competency
- Harmful Content: Suggesting dangerous treatments

# Continuous Learning Systems

## Adaptive Medical AI

Continuous learning allows RLHF systems to adapt to new medical knowledge, evolving guidelines, and emerging conditions.

## Learning Pipeline

- Data Collection: Continuously gather new preferences and outcomes
- Quality Control: Validate and filter incoming feedback
- Incremental Training: Periodically update models
- Validation: Test updated models before deployment
- Staged Rollout: Gradual deployment with monitoring

## Update Frequency

- Critical Updates: Immediate for safety issues
- Guideline Changes: When major clinical guidelines update
- Performance Improvement: Monthly or quarterly cycles
- Seasonal Adjustments: For seasonal conditions (e.g., flu)

# Feedback Loop Design

## Closed-Loop Learning

Effective feedback loops ensure continuous improvement from deployment experience back into the training process.

## Loop Components

- Deployment: Model serves real clinical interactions
- Monitoring: Log outputs, user interactions, outcomes
- Expert Review: Clinicians evaluate flagged cases
- Data Curation: Select valuable examples for training
- Model Update: Retrain with new feedback data
- Validation & Rollout: Test and deploy improved model

## Feedback Triggers

- Low Confidence: Model uncertainty about output
- User Disagreement: User corrects or rejects output
- Adverse Events: Negative outcomes reported
- Novel Cases: Situations not seen during training

# A/B Testing Framework

## Experimental Design

A/B testing allows safe evaluation of RLHF model updates by comparing new versions against existing systems.

## Test Setup

- Control Group: Existing model (baseline)
- Treatment Group: New RLHF model variant
- Randomization: Assign users/cases randomly to groups
- Sample Size: Calculate based on expected effect size
- Duration: Run until statistical significance achieved

## Evaluation Metrics

- Primary: Clinical accuracy, safety events
- Secondary: User satisfaction, efficiency, cost
- Guardrail: Metrics that must not worsen (safety)
- Exploratory: Additional insights (e.g., bias, fairness)

# Performance Monitoring

## Continuous Surveillance

Real-time monitoring detects performance degradation, safety issues, and opportunities for improvement in deployed RLHF systems.

## Key Performance Indicators (KPIs)

- Accuracy Metrics: Diagnostic accuracy, treatment appropriateness
- Safety Metrics: Error rates, adverse event reports
- Efficiency Metrics: Response time, query resolution rate
- User Metrics: Satisfaction scores, engagement rates
- Technical Metrics: Latency, uptime, system load

## Monitoring Dashboard

- Real-Time Alerts: Immediate notification of critical issues
- Trend Visualization: Charts showing metric evolution over time
- Anomaly Detection: Automated identification of unusual patterns
- Drill-Down Analysis: Investigate specific cases or time periods
- Comparative Views: Compare across versions, time periods, demographics

# Failure Mode Analysis

## Understanding Failures

Systematic analysis of failure modes helps prevent future errors and improves RLHF system robustness.

## Common Failure Modes

- Overconfidence: Model too certain about incorrect outputs
- Underspecification: Ambiguous queries lead to inappropriate responses
- Distribution Shift: Performance drops on out-of-distribution inputs
- Reward Hacking: Model exploits loopholes in reward function
- Bias Amplification: RLHF reinforces training data biases

## Root Cause Analysis

- Data Issues: Insufficient or biased training data
- Model Limitations: Architecture can't capture necessary patterns
- Reward Misspecification: Reward doesn't capture true objectives
- Training Instability: Optimization issues during RLHF
- Deployment Mismatch: Different conditions than training

# Case Study: Treatment Recommendation System

## System Overview

A real-world example of RLHF applied to an AI system that recommends treatment options for chronic disease management.

## Implementation Details

- Base Model: Fine-tuned medical LLM (e.g., Med-PaLM)
- Preference Data: 50,000 comparisons from 200 physicians
- Reward Model: Transformer-based classifier on treatment quality
- Policy Optimization: PPO with safety constraints
- Deployment: Staged rollout over 6 months

## Results

- Accuracy: 15% improvement in treatment appropriateness
- Safety: 40% reduction in contraindication errors
- Satisfaction: 85% physician approval rating
- Efficiency: 30% reduction in time to formulate treatment plan
- Adherence: 12% increase in patient treatment adherence

# Hands-On: Building an RLHF Pipeline

## Practical Implementation

Step-by-step guide to implementing a basic RLHF pipeline for medical applications.

## Setup & Prerequisites

- Python 3.8+, PyTorch, Transformers library
- Pre-trained medical language model (e.g., BioBERT, ClinicalBERT)
- Preference dataset (or use publicly available data)
- GPU with 16GB+ VRAM recommended

## Pipeline Steps

- Step 1: Load and prepare preference dataset
- Step 2: Train reward model on preferences
- Step 3: Set up PPO or DPO training loop
- Step 4: Optimize policy with reward guidance
- Step 5: Evaluate on held-out test set
- Step 6: Analyze outputs and iterate

# Ethical Considerations in Medical RLHF

## Ethical Imperatives

RLHF in healthcare must navigate complex ethical terrain, balancing innovation with patient safety and equity.

## Key Ethical Concerns

- Bias & Fairness: Ensure equitable performance across demographics
- Transparency: Make AI reasoning understandable to clinicians
- Accountability: Clear responsibility when errors occur
- Privacy: Protect patient data in feedback loops
- Autonomy: Preserve patient and physician decision-making
- Beneficence: Ensure AI improves outcomes for all patients

## Bias Mitigation

- Diverse Annotators: Include experts from various backgrounds
- Stratified Evaluation: Test performance across demographic groups
- Fairness Metrics: Monitor for disparate impact
- Bias Audits: Regular third-party fairness assessments

# Thank you

## Key Takeaways

- ✓ RLHF aligns AI with medical expertise through expert feedback
- ✓ Safety constraints are critical in healthcare applications
- ✓ Continuous monitoring ensures long-term performance
- ✓ Ethical considerations guide responsible deployment

Ho-min Park

[homin.park@ghent.ac.kr](mailto:homin.park@ghent.ac.kr)

[powersimmani@gmail.com](mailto:powersimmani@gmail.com)