

Lecture 13 - Contents

An overview of the main sections in this lecture.

Part 1

Federated Learning Overview

Part 2

Privacy and Communication Efficiency

Part 3

Case Studies and Benchmarks

Hands-on

Federated Learning Simulation

This outline is for guidance. Navigate the slides with the left/right arrow keys.

Lecture 13:

Federated Learning for Medical LLMs

Privacy-Preserving Medical AI




Federated Learning Overview

FL Principle

Collaborative learning without centralizing data.
Models travel to data, not vice versa.

Medical Privacy

HIPAA & GDPR compliant. Patient data never leaves the hospital premises.



Data stays distributed • Models aggregate learnings • Privacy preserved

Part 1

Privacy-Preserving Techniques



Encryption




Anonymization



Secure Aggregation

Distributed Architecture



 Gradient Upload 

Client-Server Communication


Communication Protocol

gRPC: High-performance RPC framework

REST API: HTTP-based communication

WebSocket: Real-time bidirectional

Message Flow



Security Measures



TLS 1.3



Mutual Auth



Rate Limiting



Audit Logging



Verification

Aggregation Algorithms



Weighted Average

$$w_{\text{global}} = \sum(n_i/N \times w_i)$$



Weight by dataset size



FedAvg

Most popular algorithm



Simple & effective




FedProx

Adds proximal term



Handles heterogeneity

Aggregation Process




FedAvg & FedProx




FedAvg


$$w = \sum(n_k/n \times w_k)$$



+



+



→



- Simple weighted average
- Fast convergence
- Works well with IID data



FedProx

$$\min F(w) + \mu/2\|w-w_t\|^2$$

Standard
Loss $F(w)$

+

Proximal Term
 $\mu/2\|w-w_t\|^2$



- Adds proximal term
- Handles heterogeneity
- More robust

Differential Privacy

Original Data

+

Calibrated Noise

=

Protected Data 

Mechanisms

Gaussian Mechanism



Laplace Mechanism



Privacy Budget (ϵ)

ϵ remaining

Parameters

ϵ

Privacy Loss
Lower = More Private

δ

Failure Probability
Typically 10^{-5}

Typical Values


$\epsilon = 1-10$
 $\delta = 10^{-5}$

Secure Aggregation

Server learns only aggregate, not individual updates

Process

- 1 Clients mask updates
 $u_1 + m_1 = w_1$ masked
- 2 Server aggregates
 $\Sigma(masked) = \Sigma(w_i)$
- 3 Masks cancel in sum
 $m_1 + m_2 + m_3 = Q$



Shamir Secret Sharing

Split secrets into shares


Homomorphic Encryption

Compute on encrypted data

Secure Multi-Party Computation

Joint computation without revealing

Homomorphic Encryption



How It Works

Mathematical Property:

$$E(3) + E(5) = E(8) \quad E(a) \oplus E(b) = E(a \oplus b)$$

Operation on encrypted data = Encryption of operated data

Types

Partial HE

Only addition OR
multiplication

Somewhat HE

Limited number of
operations

Fully HE ★

Arbitrary
computations

**Security**

100-1000x slowdown

Performance⚡

Part 2

Medical Data Challenges


Data Heterogeneity



Hospital A Hospital B Hospital C

Statistical

Different data distributions
across hospitals



System


Varying compute resources and
network



Model

Different architectures and
objectives

Non-IID Medical Data



⚠ Different patient demographics

Age, gender, ethnicity vary by location



⚠ Varying disease prevalence

Regional differences in disease patterns



⚠ Hospital-specific protocols

Different treatment guidelines and standards



⚠ Equipment differences

Imaging devices, sensors vary in quality



⚠ Regional health patterns

Climate, lifestyle, socioeconomic factors



Impact: Model bias • Slower convergence • Reduced accuracy

Client Drift Handling

FedProx

Proximal term limits drift

SCAFFOLD

Variance reduction

FedDyn

Dynamic regularization

Adaptive LR

Adjust learning rates

Communication Efficiency

Gradient Compression

Full
Gradients
100 MB




Compressed
1-10 MB

Sparsification, quantization

10-100x reduction

Local Updates



More local epochs

Linear in E reduction

Model Compression



Full Model

Prune
Distill



Compressed

Pruning, distillation

2-10x reduction

Model Personalization

- ✓ Fine-tune last layers locally
- ✓ Meta-learning (MAML)
- ✓ Mixture of global & local models
- ✓ Clustered federated learning


Global knowledge  Local adaptation

Cross-Silo vs Cross-Device



Cross-Silo (Hospital FL)

~10-100 clients



- Few clients (~10-100)
- Large datasets
- Reliable connectivity
- Powerful compute



Cross-Device (Mobile FL)

Millions of clients



- Many clients (millions)
- Small local data
- Unreliable connectivity
- Limited resources

Part 3

Multi-Hospital Collaborations



Multi-Institutional Studies

Broader Data Diversity

Access to diverse patient populations across regions and demographics

Larger Sample Size

Combined datasets enable robust statistical analysis

Privacy Preserved

Collaboration without sharing raw patient data

Better Generalization

Models trained on heterogeneous data perform better

Research Outcomes

- ✓ Faster rare disease diagnosis
- ✓ Improved treatment predictions
- ✓ Reduced model bias
- ✓ Enhanced clinical decision support

Regulatory Compliance

EU GDPR

European Union General Data Protection Regulation

- Right to explanation
- Data minimization
- Privacy by design
- Cross-border transfer rules

US HIPAA

Health Insurance Portability and Accountability Act

- PHI protection
- Security safeguards
- Breach notification
- Business associate agreements

✓ Compliance Checklist

- Encrypted communication channels
- Audit logs and monitoring
- Data anonymization techniques
- Consent management
- Regular security assessments

Data Governance



Governance Structure

Clear roles and responsibilities



Policy Framework

Data usage policies and procedures



Access Control

Role-based permissions



Audit Trail

Complete logging of data access

Quality Control

✓ Data Validation

Schema validation and quality checks

Model Testing

Comprehensive evaluation metrics



Performance Monitoring

Continuous accuracy tracking



Bias Detection

Fairness and equity assessments

Incentive Mechanisms



Financial Rewards

Payment for data contribution



Reputation System

Recognition and credibility



Data Credit

Fair attribution of contributions




Performance Bonus

Rewards for quality data


Performance Benchmarks

⚖️ Centralized vs FL



90-95% accuracy retention in FL

⌚ Training Time



2-5x slower due to communication

💾 Bandwidth Usage



1-10 GB per round

🎯 Convergence



10-50 rounds to converge

Security Auditing



Vulnerability Analysis

Regular penetration testing



Compliance Audits

GDPR/HIPAA checks



Threat Modeling

Identify attack vectors



Security Metrics

Track security posture

Scalability Analysis

Node Scaling

Linear scaling up to 100s of nodes

Bottlenecks

Network & aggregation server

Load Balancing


Distribute training rounds

Resource Management

Optimize memory & compute

Case Study: COVID-19 FL Consortium

Global Collaboration for Pandemic Response



20+

Countries



100+

Hospitals



1M+

Patient Records



Key Achievements

 Developed COVID-19 severity prediction models

 Achieved 85%+ accuracy across diverse populations

85%

 Rapid deployment during critical pandemic phases

 Fast

 Full GDPR/HIPAA compliance maintained

 GDPR  HIPAA

 Cross-border collaboration without data transfer

 No transfer

Hands-On: Federated Setup

Flower

Friendly FL framework with simple API

```
pip install flwr  
flwr.client.start()
```

PySyft

Privacy-preserving ML library

```
pip install syft  
sy.VirtualMachine()
```



Setup Steps

- 1 Install framework (Flower/PySyft)
- 2 Define model architecture
- 3 Configure server & clients
- 4 Implement training loop
- 5 Run federated training
- 6 Evaluate global model

Future Directions



Vertical FL

Different features from different institutions



Genomic FL

Privacy-preserving genetic research



LLM Fine-tuning

Federated adaptation of medical LLMs



Split Learning

Hybrid FL + split computation



Research Opportunities

Communication efficiency • Adaptive aggregation • Cross-silo heterogeneity • Byzantine robustness • Incentive mechanisms • Real-world deployments

Thank You!



Key Takeaways

- ✓ FL enables privacy-preserving multi-hospital collaboration
- ✓ Differential privacy & secure aggregation protect data
- ✓ Heterogeneity requires specialized algorithms (FedProx, SCAFFOLD)
- ✓ GDPR/HIPAA compliance is achievable with proper design
- ✓ Real-world deployments show promising results



Resources: Flower Framework • PySyft • NVIDIA FLARE



Papers: FedAvg • FedProx • Secure Aggregation