Lecture 15 - Contents

An overview of the main sections in this lecture.

This outline is for guidance. Navigate the slides with the left/right arrow keys.

Lecture 15:

Explainable Medical AI: Building Trust Through Transparency

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com

Lecture Contents

Part 1/3:

Attention-Based Interpretability

1. Attention Visualization
2. Layer-wise Relevance Propagation
3. Gradient-Based Attribution
4. Integrated Gradients
5. SHAP for Medical Applications
6. LIME for Clinical Text
7. Concept Activation Vectors

Attention Visualization

Medical Image (X-ray)

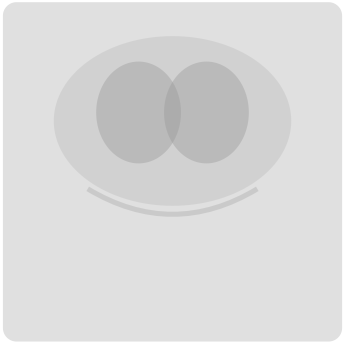

Attention Heatmap

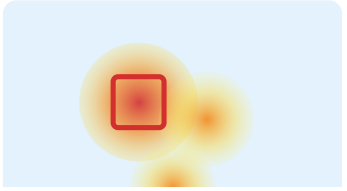

Attention Mechanism

Query (Q) | Key (K) | Value (V

Attention Score

$$\text{Attention}(Q,K,V) = \text{softmax}(QK^T / \sqrt{d_k}) \times V$$

↓

Weighted Features
Focus on relevant regions

Attention Intensity

High | Medium | Low | No

Attention Mechanism

Neural network component showing where the model focuses when making predictions

Heatmap Visualization

Color-coded maps indicating attention weights across different input regions

Medical Image Analysis

Highlighting relevant regions in X-rays, CT scans, and MRI images

Clinical Application

Identifying pathological features automatically and visualizing model focus areas

Key Benefit

Allows clinicians to verify that AI models are focusing on medically relevant features

Layer-wise Relevance Propagation (LRP)

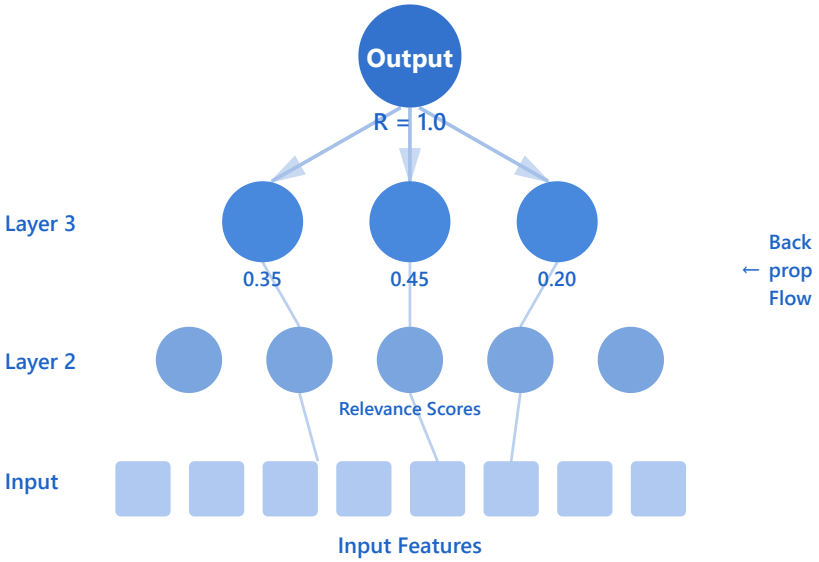

LRP Principle

Backpropagating relevance scores from output to input layers

Layer Contribution

Measuring each layer's contribution to the final prediction

Conservation Property

Total relevance remains constant through all layers

Medical Imaging

Identifying diagnostic features in medical images

Gradient-Based Attribution

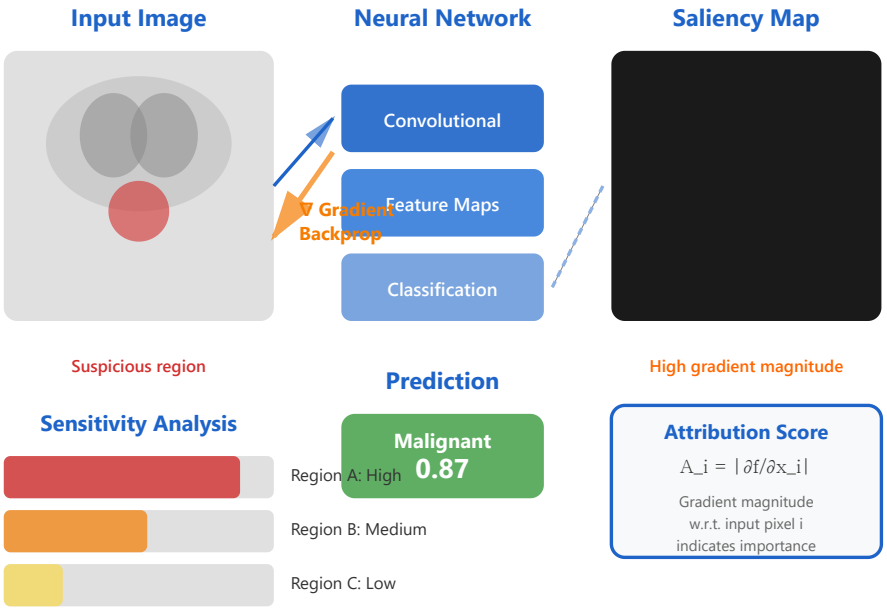

📈 Gradient Analysis

Using gradients to measure input feature importance

🎯 Sensitivity Mapping

Showing which inputs affect the output most significantly

🗺️ Saliency Maps

Visualizing important pixels and regions in images

🏥 Clinical Interpretation

Understanding what factors drive model decisions

Integrated Gradients

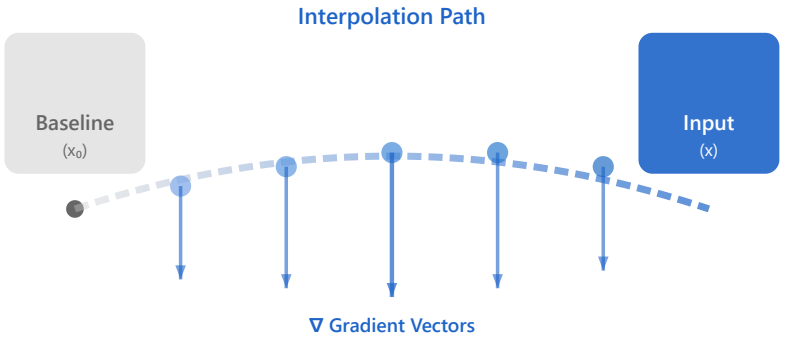

Attribution Formula

$$IG(x) = (x - x_0) \times \int_{\alpha=0}^{\alpha=1} \nabla F(x_0 + \alpha(x - x_0))\, d\alpha$$

Sum of gradients along path

∫ Path Integration

Integrating gradients along interpolation path from baseline

🔍 Baseline Comparison

Comparing with neutral reference input (e.g., black image)

✓ Axiom Satisfaction

Meeting sensitivity and implementation invariance axioms

💪 Robust Attribution

More stable and reliable than simple gradient methods

SHAP Values for Medical Predictions

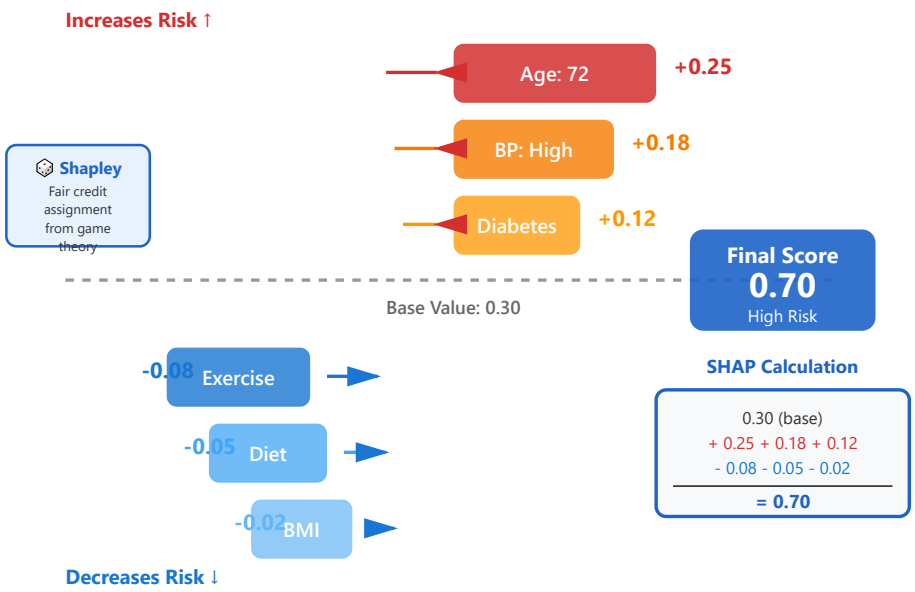

Shapley Values

Game-theoretic approach to measuring feature importance

Additive Explanations

Consistent and locally accurate explanations

Feature Contribution

Quantifying each feature's impact on the prediction

Clinical Decision Support

Explaining risk scores, diagnoses, and treatment recommendations

LIME for Clinical Text

🔍 Local Approximation

Explaining individual predictions with simple interpretable models

🔀 Text Perturbation

Systematically modifying input text to test importance

📝 Word Importance

Highlighting influential words and phrases in clinical notes

📋 EHR Text Analysis

Explaining decisions on electronic health record text data

Concept Activation Vectors (CAV)

💡 High-level Concepts

Testing for abstract medical concepts in neural networks

➡️ Direction Vectors

Finding concept directions in model activation space

🧪 Sensitivity Testing

Measuring model response to specific medical concepts

🔬 Medical Concepts

Detecting learned patterns like 'inflammation' or 'tumor characteristics'

Part 2/3:

Clinical Explanation Generation

1. Natural Language Explanations
2. Counterfactual Generation
3. Decision Path Visualization
4. Uncertainty Communication
5. Evidence Highlighting
6. Contrastive Explanations

Natural Language Explanations

💬 Text Generation

Generating human-readable explanations for predictions

🗣️ Clinical Language

Using medical terminology appropriate for healthcare professionals

🖇️ Multi-modal Integration

Combining image analysis with textual explanations

📖 Explanation Quality

Ensuring accuracy, completeness, and clinical relevance

Counterfactual Explanations

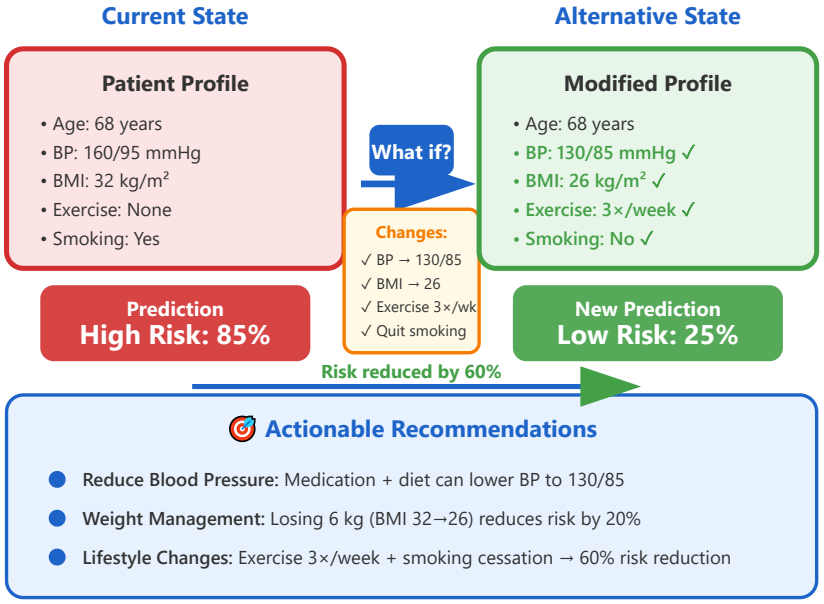

🔄 What-If Scenarios

Showing minimal changes needed to alter the prediction

🎯 Actionable Insights

Identifying modifiable factors that influence outcomes

📊 Clinical Utility

Helping clinicians understand treatment alternatives

⚖️ Minimal Changes

Finding smallest modifications for desired outcome changes

Decision Path Visualization

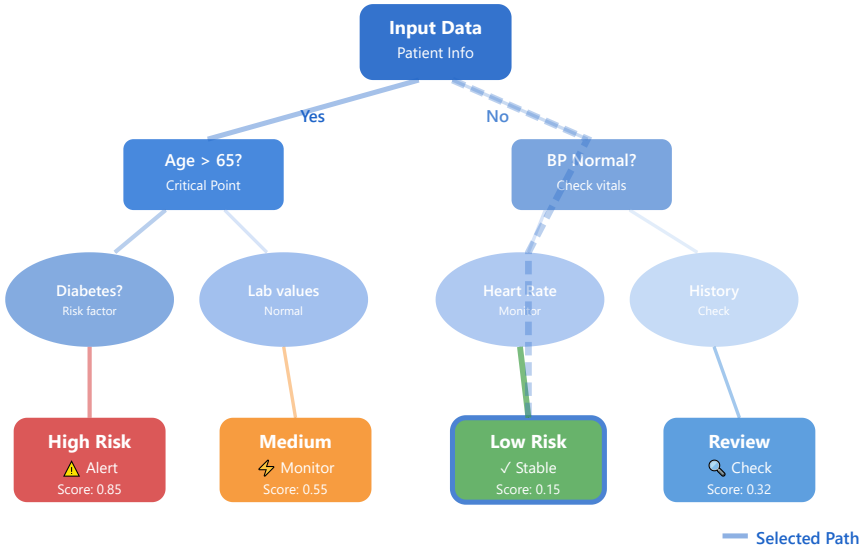

Decision Trees

Visualizing hierarchical decision-making processes

Path Tracing

Following the model's reasoning from input to output

Critical Points

Identifying key decision points in the prediction process

Visual Navigation

Interactive exploration of model decision logic

Uncertainty Communication

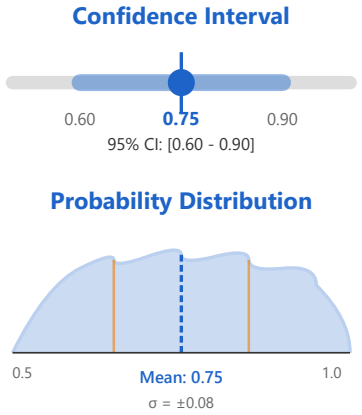

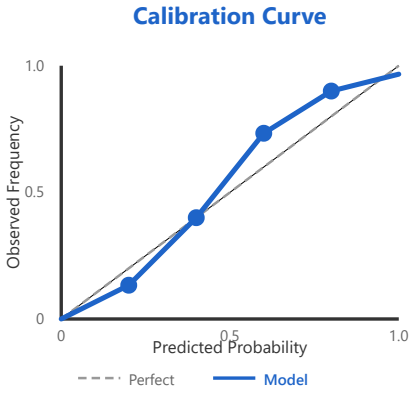

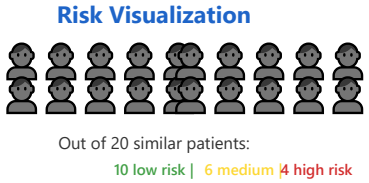

📊 Confidence Intervals

Quantifying prediction uncertainty with statistical bounds

🎲 Probabilistic Outputs

Expressing predictions as probability distributions

⚠️ Risk Communication

Clearly conveying model confidence to clinicians

🎯 Calibrated Predictions

Ensuring predicted probabilities match actual frequencies

Evidence Highlighting

✨ Feature Marking

Highlighting relevant features that support the prediction

🔗 Source Attribution

Linking predictions to specific data sources

📑 Evidence Ranking

Ordering supporting evidence by importance

🔍 Traceability

Enabling verification of model reasoning process

Contrastive Explanations

⚖️ Comparison Analysis

Explaining why A rather than B was predicted

🔄 Differential Features

Identifying key differences between alternative outcomes

🎯 Foil Cases

Using contrasting examples to clarify decisions

💡 Enhanced Understanding

Improving comprehension through comparative reasoning

Part 3/3:

Clinical User Requirements

1. Physician Interpretability Needs
2. Patient-Facing Explanations
3. Regulatory Documentation
4. Audit Trail Generation
5. Trust Calibration
6. Error Analysis & Debugging

Physician Interpretability Needs

👩‍⚕️ Clinical Workflow

Integrating explanations into existing medical workflows

⏱️ Time Efficiency

Providing quick, actionable insights without overwhelming detail

🎓 Medical Expertise

Matching explanation complexity to physician knowledge level

✅ Trust Building

Fostering appropriate trust through transparent AI

Patient-Facing Explanations

Accessible Language

Using simple, non-technical terms for patient comprehension

Visual Communication

Employing graphics and diagrams for better understanding

Empathy & Support

Providing explanations with emotional sensitivity

Health Literacy

Adapting to varying levels of medical knowledge

Regulatory Documentation Requirements

📋 FDA Compliance

Meeting regulatory requirements for AI/ML medical devices

📝 Documentation Standards

Maintaining comprehensive model development records

🔍 Validation Evidence

Providing interpretability as part of validation process

✓ Approval Process

Supporting regulatory submissions with explainability data

Audit Trail Generation

📊 Decision Logging

Recording all model inputs, outputs, and reasoning steps

🕐 Temporal Tracking

Timestamping predictions and model versions

🔒 Immutable Records

Creating tamper-proof logs for legal compliance

🔍 Retrospective Analysis

Enabling investigation of past predictions

Trust Calibration

⚖️ Appropriate Trust

Balancing between over-reliance and under-utilization

📈 Performance Awareness

Communicating model strengths and limitations clearly

⚠️ Failure Cases

Highlighting scenarios where model may be unreliable

🎯 Confidence Alignment

Ensuring user trust matches actual model performance

Error Analysis Tools

🔍 Failure Detection

Identifying systematic errors and edge cases

📊 Error Classification

Categorizing different types of prediction failures

🎯 Root Cause Analysis

Understanding why specific errors occur

🔧 Improvement Insights

Using error patterns to guide model refinement

Debugging Interfaces for Developers

Interactive Tools

Visual interfaces for exploring model behavior

Layer Inspection

Examining activations and weights at each layer

Performance Metrics

Real-time monitoring of model performance indicators

Bug Identification

Tools for detecting and fixing model issues

Performance vs. Interpretability Trade-off

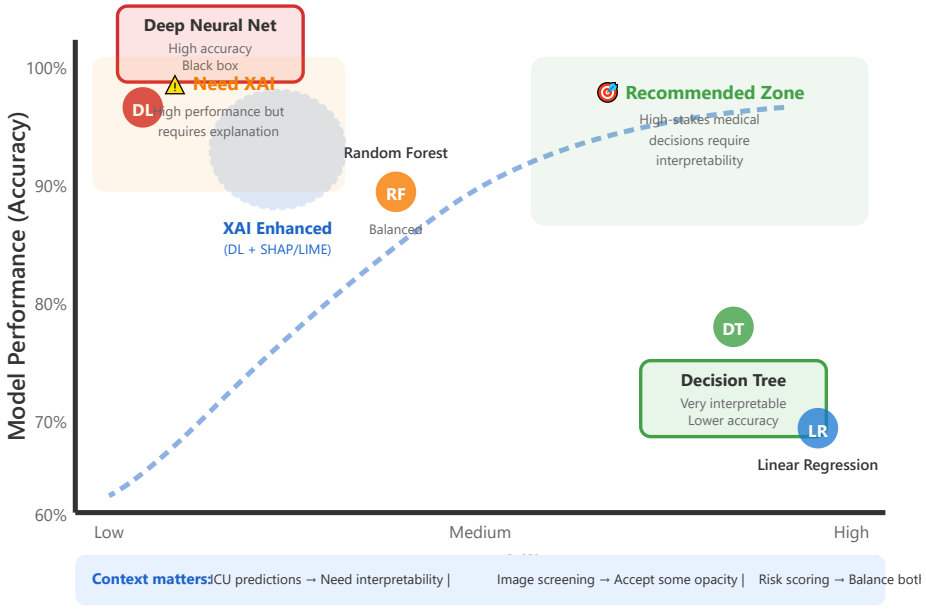

⚖️ Accuracy vs. Transparency

Balancing model performance with explainability

🧩 Model Complexity

Simpler models are more interpretable but may be less accurate

🎯 Application Context

High-stakes medical decisions may require more interpretability

💡 Hybrid Approaches

Combining powerful models with post-hoc explanations

Case Study: ICU Mortality Predictions

🏥 Clinical Context

Predicting patient mortality risk in intensive care units

📊 Feature Importance

Identifying vital signs and lab values driving predictions

⏱️ Temporal Patterns

Explaining how patient trajectory affects risk scores

👨‍⚕️ Clinical Validation

Physicians reviewing and validating AI explanations

Hands-On: XAI Tools Implementation

🔧 SHAP Library

Python implementation: shap.TreeExplainer, shap.DeepExplainer

🎨 LIME Package

Text and image explainers: lime.lime_text, lime.lime_image

📊 Captum (PyTorch)

Integrated Gradients, GradCAM, and other attribution methods

💻 Practice Exercise

Implementing explanations for medical image classification

Future Research Directions in XAI

🔮 Multimodal Explanations

Combining imaging, text, and structured data explanations

🤖 Foundation Model XAI

Explaining large language and vision models in medicine

⚖️ Standardization

Developing common evaluation metrics for explainability

🏥 Clinical Integration

Seamlessly embedding XAI into electronic health records

Thank you

Ho-min Park

homin.park@ghent.ac.kr

powersimmani@gmail.com