

Hands-on: Evaluation Implementation

1

Load Benchmark Dataset

→ MedQA, PubMedQA, or USMLE-style questions



2

Run Model Inference

→ Generate predictions with confidence scores



3

Calculate Metrics

→ Accuracy, F1, ECE, hallucination rate



4

Visualize Results

→ Confusion matrices, calibration plots, error analysis



5

Generate Report

→ Comprehensive evaluation summary

Recommended Tools

[HuggingFace Evaluate](#)[TorchMetrics](#)[scikit-learn](#)[Weights & Biases](#)