

PPO in Medical Applications

Proximal Policy Optimization (PPO)

PPO is the most widely used algorithm for RLHF, providing stable updates that prevent catastrophic policy changes in medical AI systems.

Key Features



Clipped Objective

Limits policy updates to prevent drastic changes



Trust Region

Maintains proximity to previous policy



Stability

Reduces training instability vs vanilla PG



Sample Efficiency

Reuses data multiple times per iteration

Training Process (Iterative Loop)

1

2

3