


# Concept Activation Vectors (CAV)



## 💡 High-level Concepts

CAVs test whether neural networks learn human-interpretable concepts

## ➡️ Direction Vectors

Concepts represented as directions in activation space

## 🧪 Sensitivity Testing

Measure how much a concept influences model predictions

## 🔬 Medical Validation

Verify if models use clinically relevant concepts