

Kubernetes for Medical Deployments

Kubernetes (K8s) orchestrates containerized ML workloads across clusters

⚙️ Core Components



Pods

Smallest unit containing containers



Services

Expose pods with load balancing



Deployments

Manage replicas & rolling updates

🔧 Configuration



ConfigMaps

Store configuration data



Secrets

Manage sensitive credentials



Volumes

Persistent data storage

📈 Scaling Features



Horizontal Pod Autoscaling (HPA)



Resource requests/limits



Node affinity for GPU scheduling