

Sparse Activation: Computational Efficiency through Selective Expert Activation

Dense Model (All Active)



All 8 Experts Active (100%)

High computation & memory usage

Sparse MoE (Top-2 Active)



Only 2 Experts Active (25%)

75% computation reduction!

Active Experts	8 / 8 (100%)
Computation	Full (100%)
Memory Usage	High
Latency	High

Active Experts	2 / 8 (25%)
Computation	75% Reduction
Memory Usage	Low
Latency	Low

⚡ Efficiency Gains

Sparse activation enables models with billions of parameters to run with the computational cost of much smaller models. A 64-expert MoE with Top-2 routing achieves **32x parameter scaling with only 2x computation increase!**