

## Extended Context Windows: 100K+ Tokens

### Traditional Limits

- BERT: 512 tokens
- GPT-3: 2,048 tokens
- Limited patient history view

### Modern Long Context

- GPT-4 Turbo: 128K tokens
- Claude 3: 200K tokens
- Gemini 1.5: 1M tokens

### Memory Requirements

- Attention memory:  $O(n^2)$
- 100K context: ~40GB VRAM
- Optimization crucial for deployment

### Medical Applications

- Full patient history analysis
- Multi-visit trend detection
- Comprehensive literature review

### Context Scaling

Extended context windows enable processing entire patient histories spanning years of medical records, lab results, and imaging reports in a single inference.