


Flash Attention: IO-Aware Optimization



 Revolutionary Efficiency

Flash Attention achieves exact attention with $O(n)$ memory and 3-10x speedup by optimizing GPU memory access patterns, making long-context medical AI practical. **The bottleneck is I/O, not computation!**